



ROYAL INSTITUTE
OF TECHNOLOGY

School of Computer Science and Communication
Department of Theoretical Computer Science

LAB F

Firewall: Packet filtering with iptables

NAME	KTH USERNAME

DATE :: ____ - ____ - ____

TEACHING ASSISTANT'S NAME :: _____

LAB F PASSED (TA'S SIGNATURE) :: _____

Computer Security

DD2395 / HT2015

Contents

1	Introduction	1
1.1	References	1
1.2	Preparation Questions	1
2	Getting Started	4
2.1	Connecting to the Virtual Hosts	4
2.2	Naming	5
2.3	Getting Root Access	5
2.4	Setting up Interfaces	6
2.5	Setting up Routing	6
3	iptables: Building a Firewall	7
3.1	Ping and the Internet Control Message Protocol (ICMP)	7
3.2	Logging and Limits	8
3.3	Building a firewall	9
4	nmap: Detecting Server Capabilities	14
4.1	Enumeration	14
4.2	Service Discovery	14
4.3	OS discovery	16
5	Cleanup	16
6	History	16
A	Lab Network Map	17
B	FAQ – Common Problems and Fixes	17
C	Flow Diagram of iptables	18
D	Command Line Basics – in case you are not familiar with the shell	19
D.1	Manual Pager	19
D.2	Viewing a File	19
D.3	Editing a File	19
D.4	Background Jobs	19
D.5	Clipboard – Copy+Paste	19

1 Introduction

During this lab you will be introduced to the concepts of firewalls, i.e., packet filtering in GNU/Linux using [iptables](#) to set up and configure filtering rules. Moreover, you will learn how to use the network mapper tool [nmap](#) to explore systems on the Internet. The lab is also a refresher of your basic GNU/Linux and computer networking skills.

iptables (also known as netfilter) is the user space command line program used to configure the Linux 2.4.x and later packet filtering ruleset, a successor of ipchains. Both netfilter and ipchains were started in 1999 by Paul Russell, a free software developer.



Deadlines: Preparation questions before coming to the lab, tasks by the end of the session

The questions in Section 1.2 are to be answered **before** coming to the lab session you signed up for. The mandatory part of the lab (Section 2 and Section 3) is to be finished by the end of the lab session you signed up for. Therefore, you are encouraged to come prepared and use the time wisely.

1.1 References

Before and during the lab, you might want to consult some of the references that you will find below.

For iptables, a concise overview of the most important commands can be found at:

<https://www.centos.org/docs/4/html/rhel-rg-en-4/s1-iptables-options.html>

Two more extensive documentations, including a refresher of general networking knowledge, can be found at:

<https://www.frozentux.net/documents/iptables-tutorial/>

<https://www.netfilter.org/documentation/HOWTO//packet-filtering-HOWTO.html>

For nmap there is a reference at: <https://nmap.org/book/man.html>

Especially when looking for the syntax and the meaning of arguments, the manual pages are a good reference (on the command line type `man nmap` or `man iptables`, to search type `/searchterm` and press enter, to repeat the search type `n`, to quit press `q`).

If you are not used to working on the command line, have a look at Appendix D.

1.2 Preparation Questions

Please answer the following questions before coming to the lab session. All questions can be answered with the references mentioned in Section 1.1.

1.2.1 iptables

Chains The rules in iptables are always parts of a chain. Chains can either be user created or one of the built-in chains. In this lab we will use three standard chains: INPUT, FORWARD and OUTPUT. Explain which packets will pass through each of these chains:

INPUT: _____

FORWARD: _____

OUTPUT: _____

Operators To change a chain, you need to use an operator. Look up the short form command for the following operators and explain briefly what they do:

Append: _____

Insert: _____

Delete: _____

List: _____

Flush: _____

Filters Filters are used to choose which packets match a rule. Each filter is used to create matches. Some matches requires modules to be loaded with `-m MODULENAME` to be available. Explain the following matches and how they can be used:

-p: _____

-s: _____

-d: _____

-i: _____

-o: _____

--sport: _____

--dport: _____

Jump Targets Jump targets are used to decide what to do with a packet once it has matched a rule. Explain the following targets:

ACCEPT: _____

DROP: _____

REJECT: _____

LOG: _____

1.2.2 nmap

The nmap tool provides different scan methods, to discover and analyse remote hosts. Each of these scan methods uses a certain network protocol. Each network protocol operates on a certain [OSI-model](#) layer (e. g., ARP on OSI layer 2, IP on OSI layer 3, etc.). So we can group the nmap scan methods according to OSI layers. Please name some examples for each group mentioned below and describe them (name of the scan method, what command/arguments starts it, which network protocol does the scan method use, what is the purpose of the scan method, what

are the limitations of the scan method).

OSI layer 3:

OSI layer 4:

OSI layer 7:

**Milestone**

Report your progress to a lab assistant. _____

2 Getting Started

Saving time: Work in parallel and ask questions (using “Stay A While”)

In order to be able to complete the lab during the allocated time, you can do more than one task at a time as not all of them have to be done sequentially. If you have to wait for a program to finish or a milestone to be signed, consider continuing with another part and/or start reading the details. If you get stuck somewhere, don't hesitate to ask for help (using the queuing system “Stay A While”, choosing the Dasak queue: <http://queue.csc.kth.se/#/queue/Dasak>).

The laboratory system consists of a virtualized system with three virtual hosts that simulate two interconnected networks (see Appendix A). These hosts will be running GNU/Linux (Ubuntu Server 14.04.3).

You will get a user account on the xen machine (the hypervisor machine running the virtual hosts) and logins to the virtual hosts at the beginning of the lab session. The IP addresses for your network are listed in the network topology map in Appendix A. Please, note that the virtual hosts will not have access to the general Internet.

Group Number, Letter and Login Credentials

The following information will be handed out to you at the beginning of the lab session.

group number #: _____ (used in the IP addresses)
group letter □: _____ (used in the username and host names)

xen login (xen.netlab.csc.kth.se)

username: student-□
passphrase: _____

virtual hosts (iptables-□.1, iptables-□.2 and iptables-□.3)

username: student
passphrase: basket field payment same love
(words are separated with spaces¹)

2.1 Connecting to the Virtual Hosts

Login to the virtualization host `xen.netlab.csc.kth.se` via `ssh` using the `student-□` user with the passphrase supplied to you:

```
ssh student-□@xen.netlab.csc.kth.se
```

You will end up in the `virt-shell` environment where you can manage the virtual machines you have been assigned to. There, boot and login to one of the three virtual machines `iptables-□.1` (internal), `iptables-□.2` (firewall), or `iptables-□.3` (external). You can login by first choosing the instance, for example:

¹This passphrase was generated using the diceware method (see <http://world.std.com/~reinhold/diceware.html>). When used correctly, it yields secure passphrases that are easier to memorize. This passphrase was generated by randomly selecting five out of 850 words and therefore has an entropy of 48.7 bits (5 · 9.73 bits). This is the same security level that you would get with a password composed of 8 random characters (using lower and upper letters, numbers and eight special characters: 8 · 6.13 bits = 49.0 bits of entropy) which for most people would be harder to memorize. Note that for passphrases that have to protect sensitive information, you would probably want to have a level of about 90-100 bits of entropy, so for example seven or eight words randomly selected from the full diceware wordlist (consisting of 7776 short words).

```
control iptables-A.1
```

booting it (if it is not running):

```
boot
```

and then use the

```
console
```

command in order to proceed to the login shell (you might need to wait a few seconds and then press Enter before you see the login prompt).

Repeat these steps for each of the three virtual hosts in a new terminal window.



Working with the Virtual Serial Consoles

- Use 3 terminals for the 3 different virtual hosts.
- Do not connect from multiple consoles to the same virtual host. This will mess up your consoles.
- For line-wraps to work correctly the size of your terminal window has to match the virtual console. If you resized your terminal window (e. g., maximized it), you need to issue the command `resize` to make line-wraps work again. If that does not work, you can either make sure your terminal window has the default size (80×24) or update the virtual console window size (using `stty columns X rows Y`). The command `stty -a` prints the current settings and lets you get the size of a window not connected on a serial console.
- If your virtual console seems strange, try the `clear` command. It will try to reset the console.

2.2 Naming

To ensure you are on the right machine, set the **hostname** as soon as you have logged in, and then log out and in again. For example:

```
student@iptables:~$ sudo hostname iptables-A1
student@iptables:~$ logout
Ubuntu 14.04 iptables-A1 ttyS0
iptables-A1 login: student
Password:
...
student@iptables-A1:~$
```

Do this on all three virtual hosts. You may get “unable to resolve host” error messages after issuing the `hostname` command, which you can safely ignore (the hosts don’t have Internet access and therefore fail to do a DNS lookup).

2.3 Getting Root Access

The next step is to permanently get root access on the virtual hosts.

```
student@iptables-A1:~$ sudo bash
root@iptables-A1:~#
```

Here the `#` indicates that you have root access on the host. While it is a bad habit (normally you should use `sudo` instead of a root shell), almost all commands in this lab requires root access, so this will save you from typing `sudo` constantly. Do this on all three virtual hosts.

2.4 Setting up Interfaces

Next, you need to configure your network interfaces. You do this by manually assigning them an IP address and a netmask. The network map in the appendix describes which addresses to use. Remember to replace # with your group number in all IP addresses (for example if your # is 4 and your □ is E, then the host `iptables-E.1` has the IP address 192.168.4.2). Use the `ifconfig` tool and make sure the configuration is applied on the correct interfaces (`eth0`, `eth1`) on each machine.

For example, to assign the IP address 10.18.0.33 with a 16 bit netmask to the interface `eth0`, issue the following command:

```
ifconfig eth0 10.18.0.33/16
```

Check the configuration by using `ifconfig` without any arguments.

2.5 Setting up Routing

At this point it should be possible to ping between directly connected hosts (`iptables-□.1 ↔ iptables-□.2` and `iptables-□.2 ↔ iptables-□.3`) but *not* between the two end hosts (`iptables-□.1 ↔ iptables-□.3`). Verify this, using the `ping` command.

To ensure that data can be routed through the firewall host (`iptables-□.2`), you have to setup forwarding. First, verify that packet forwarding is enabled on the firewall host.

This is done by checking, that the file `/proc/sys/net/ipv4/ip_forward` contains “1”:

```
cat /proc/sys/net/ipv4/ip_forward
```

If it is not, set it to 1 to enable forwarding (only on the firewall host `iptables-□.2`):

```
echo 1 > /proc/sys/net/ipv4/ip_forward
```

Then, only on the end hosts `iptables-□.1` (internal) and `iptables-□.3` (external) use the command

```
route add default gw ADDRESS
```

to set their default rout to the central firewall host `iptables-□.2`. ADDRESS has to be the corresponding address of the directly connected interface on the firewall host.

Do not forget to check the routing table with the `route` command without arguments (be patient, the output may be delayed a bit).



Milestone

Verify that you can ping between the two end hosts of your network (`iptables-□.1 ↔ iptables-□.3`).



3 iptables: Building a Firewall

Note that all the `iptables` commands have to be run **only on the firewall host** (iptables-□.2)!

To view the `iptables` help, run the command:

```
man iptables
```

More references are listed in Section 1.1.

To list the current state (list of rules and policies for all chains), run the command:

```
iptables -vL
```



Saving your work

In case a virtual host fails, you might lose the rules you have created. Therefore keep a record of what you are doing (on paper or in a file on your local machine). At any time you can use the `iptables-save` command, to print out the rules you created so far. Note that saving them to a file on the virtual host might not be a good idea, as the complete host can fail.

3.1 Ping and the Internet Control Message Protocol (ICMP)

The `ping` tool is usually used to test the reachability of a host that implements the Internet protocol (IP). It operates by sending packets using the Internet Control Message Protocol (ICMP) of the type ‘echo request’ to the host whose reachability is to be tested, and processing the response from that host (if any).

In Listing 1, you can see a sample output of this utility testing the web server of KTH (`ping www.kth.se`). The results show three ICMP requests (which were successful as the time implies that there was a response) and a statistical summary of the response packets received.

```
1 PING www.kth.se (130.237.32.143) 56(84) bytes of data.  
64 bytes from 130.237.32.143: icmp_req=1 ttl=254 time=10.4 ms  
3 64 bytes from 130.237.32.143: icmp_req=2 ttl=254 time=10.4 ms  
64 bytes from 130.237.32.143: icmp_req=3 ttl=254 time=9.31 ms  
5 --- www.kth.se ping statistics ---  
3 packets transmitted, 3 received, 0% packet loss, time 2002ms  
7 rtt min/avg/max/mdev = 9.315/10.064/10.478/0.543 ms
```

Listing 1: Sample output in the command line of the `ping` utility

3.1.1 Blocking ICMP requests

`ping` the internal host iptables-□.1 from the external host iptables-□.3. What happens? Now execute the following command on the firewall host:

```
iptables -A FORWARD -p icmp --icmp-type echo-request -j DROP
```

`ping` the host again. What happens and why?

Check how the output from `iptables -vL` changed.

To remove the rule from the chain, first list the rules again:

```
iptables --line-numbers -L FORWARD
```

There you can find the line number of the rule you just added, and then remove it with:

```
iptables -D FORWARD line_number
```

3.1.2 Rejecting ICMP Requests

Now, create a new rule which **REJECTS** all ICMP echo-requests from the external network to your internal network. Note that iptables rules only match if all the conditions in the rule are true.



Specifying Source and Destination

For simplicity, we have only one internal and one external host in the lab network setup. However, when specifying the source and destination for iptables rules, try to be more general than specifying only single IP addresses as source or destination. You can accomplish this either by using an IP address range (e. g., 10.#.0.0/16) for the `-s` and `-d` option, or by specifying the involved interface with `-i` and `-o`.

Now, verify that

- you can ping from the internal host to the external,
- the external host cannot ping the internal, and
- the firewall can ping both hosts.

Can you ping from the external host to the **internal** interface `eth0` on the firewall? Why/why not?

Can this have any security implications?

What is the difference between rejecting and dropping traffic?

What are the advantages of rejecting and what are the advantages of dropping?

3.2 Logging and Limits

One important task for a firewall is to log rejected or dropped packets, making it easier to trace attacks. A rule can be created with the jump target `LOG` to save information to the system log. iptables logs directly to the syslog system.

Create a rule that logs all rejected ping messages. Make sure the string “Ping rejected by <your name>:” is written in the log message. Check the system log so the traffic is saved. This can easily be done with:

```
tail -f /var/log/syslog
```

The module `limit` can be used to limit how often a rule can be triggered. Use the `limit` module to make sure no more than 5 pings each minute are saved to the system log.



Milestone

Report your progress to a lab assistant. _____

3.3 Building a firewall

You will now build a simple firewall for your network. Before you start, make sure that there are no rules left from the previous assignments: flush all rules by running

```
iptables -F
```

(not specifying a chain removes all rules from all chains).



Order of rules

For each step, think about the order of the rules: Is it enough to append a new rule at the end of a chain, or should it go before a certain other rule, in order to work as intended? If you use the `-I` argument without specifying a number, the rule will be placed at the top of the chain.

3.3.1 Default Policy

Each chain has a default policy target that details what to do with any packet that does not match any rules present in the chain. Set all chains to the policy target **DROP**.

Verify that you can not send any data through or to the firewall (e. g., using `ping`).

Default policy targets can only be set to **ACCEPT** or **DROP** but not to **REJECT** as it is an extension target. If you would want to **REJECT** all packages that do not match a rule in the **INPUT** chain, how would you do it?

3.3.2 Network Permissions

Now it is time to start allowing some carefully chosen traffic through the firewall.

- Create a rule that allows all traffic originating from the internal network arriving to the internal interface (`eth0`) of the firewall host.
- Create a rule that allows all traffic originating from the firewall to reach the internal network.

Make sure you now can reach the internal network from the firewall, and the firewall from the internal network.

3.3.3 Permitting a Service

SSH (Secure shell) is a common service for remote management of firewalls. Create rules which allow a host from the external network to connect to the firewall with SSH. SSH runs on port 22 and uses only TCP.

Make sure that the ssh daemon `sshd` is running on all hosts by executing `service ssh restart`



Source and destination ports

Note that the well-known port number for a service only applies for the machine where the service is running (server), not for machines that connect to the service (clients). So if a client machine sends a packet to a specific service on a server machine, the destination port will be the well-known port number for that service but the source port can be any number. Accordingly, when the server machine sends an answer packet from the service back to the client machine, this packet will have the well-known port number as source port but the destination port will be the number chosen by the client machine.

Verify that,

- you **can** connect to the firewall with SSH from the other hosts,
- you **cannot** connect directly from your external host to the internal host with ssh, and

- you **can** connect from your external host to the internal host if you first connect to the firewall with ssh and then connect from there with ssh to the internal host.

What kind of security advantage does a setup with an SSH terminal server offer?

What kind of security disadvantage does a setup with an SSH terminal server introduce?

3.3.4 Stateful Filtering

In most cases we want to allow hosts on the internal network to connect to the external network, e. g., the Internet. However, we do not want hosts on the Internet to be able to connect to hosts on the internal network. For some protocols, such as TCP this can be done *stateless*, due to the three way handshake needed to create a connection: By blocking the initial SYN packet in one direction we can prevent the establishment of connections in that direction while still allowing connections that were established in the other direction to send packets both ways (as they will never send an initial SYN packet in the direction that we blocked this).

However, stateless filtering breaks a large number of protocols, for example UDP based protocols can not easily be allowed through in only one direction. Therefore, we need a *stateful* firewall to properly handle it. Some protocols, such as FTP, also break if connections from the Internet are completely denied.

Examine the module `state`. Use this module to create a match that allows the hosts on the inside of the firewall to establish connections to the outside. Allow all data that belongs to these connections through the firewall. Keep blocking all other connection attempts from the outside.

3.3.5 Opening Ports

Sometimes, you want computers on the outside of the network to have access to a specific service on the inside. Your task now is therefore to add rules to your firewall, so that external computers on the outside can reach the *echo* service (port 7) on hosts on the inside both on UDP and TCP.

The echo service (provided by the extended internet service daemon `xinetd`) is a legacy testing protocol that simply replies with the same data that is sent to it. Usually it is disabled by default, but on your hosts it is activated (you can check this by running `lsuf -i` and look for lines that end in “echo (LISTEN)” or similar). To test your firewall rules, connect to the echo service that is running on your internal host, **from the external host** by running `telnet IPADDRESS 7`

on the external host and observe that all data you send is echoed back to you. To exit telnet on the external host, you cannot use the default escape character `Ctrl+]` (as this exits to the virtual shell). Instead you have to terminate the connection on the other end: on the **internal** host kill the `xinetd` service by running

```
killall -KILL xinetd
and restart it afterwards with
service xinetd start
```

3.3.6 Blocking Ports

Sometimes you do not want your *internal* users to be able to connect to the Internet on a specific port at all. One commonly blocked port is 135 which is used for Windows file sharing². Make sure your firewall blocks all traffic on port 135 (both TCP and UDP) from the computers on the internal network.

²Actually there are more ports involved (135-139 on older Windows machines and 445 on more recent ones) but in this lab it is enough to block 135 as one example.

3.3.7 Verifying Your Setup

Verify your setup in a systematic manner, making sure all rules work correctly. To test a certain rule, you can listen to a port with `netcat` and try to connect to it using `telnet`:

On the receiving host run: `nc -l PORT`

On the sending host run: `telnet IPADDRESS PORT`

Furthermore, use the counters in the output of `iptables -vL` to make sure the packets match the rules you expect. At this point you should have the following rules active, in this order of priority:

- Connections on **port 135** from the inside hosts are blocked.
- Connections on the **SSH port** are allowed to the firewall host from the outside.
- Connections on the **echo port** are allowed to the internal hosts from the outside.
- Connections **directly to and from the firewall** are allowed from the internal networks.
- Connections **from the inside** are allowed out.
- Connections **from the outside** are blocked.



Milestone

Report your progress to a lab assistant. _____

3.3.8 Defending Against SSH Brute-force Attacks

You have allowed connections to the SSH port (hence, its service) of the firewall from any host on the outside network. There is nothing that would avoid a brute-force attack or a Distributed Denial of Service (DDoS) attack because there are no limitations on the amount of times that one particular host can try a username and a password (besides the ones that the SSH server implementation might have internally). So you are going to put in practice the knowledge you have acquired about stateful filtering and logging to prevent a host from the outside network to execute a brute force attack.

Let's start by creating a new chain called `SSH` that is supposed to handle all SSH traffic. After creating the chain, you need to modify the rule in the `INPUT` chain that allowed connections to the SSH port on the firewall host from the outside network: send all SSH packages to the '`SSH`' chain instead of directly accepting them.

In the '`SSH`' chain, first add a rule that accepts all packages from established connections, so that all the following rules only apply for new connection attempts (and don't slow down legitimate, established connections).

Now, have a look at the documentation of the module `recent`. It allows you to create a dynamic list of IP addresses for later use, such as counting the amount of times a particular IP address tries to start a new connection on a port. Your next step, with the knowledge of the module `recent`, is to create the necessary rules in the '`SSH`' chain to catch those hosts that are potentially dangerous. We will assume that anyone trying to log in **more than three times in a row within 30 seconds** from the same host is potentially malicious and should be slowed down.

We want not only to block, but also to log suspicious packages. But in case of a brute-force attack, we don't want the logfile to be flooded. So from each suspicious IP address we will block the fourth, fifth and all further connection attempts (in 30 seconds), but only log the fourth connection attempt.

Create a first rule so that the source address of the packet is added to a custom dynamic list named `SSH_COUNTER`. Then, add a second rule so that the fifth (and any further) packet attempting to establish an SSH connection and coming from the same source in a time period of 30 seconds gets dropped.

In order to log suspicious connection attempts, create two rules so that one of them logs the fourth SSH connection from the same host address (use the prefix "`SSH brute force attacker:` "), and the other one blocks the package. Note that in order to avoid getting our log full of entries for every malicious attacker we implemented the previous

rule such that the fifth attempt and any other subsequent one within the specified time gets dropped directly without being logged (and only the fourth attempt will be logged and blocked).

Finally, you have to accept all packets that have not been stopped so far. Why do you need to do this?

At this point the newly created ‘SSH’ chain should contain the following rules, in this order of priority:

- A rule accepting all packets from already established connections.
- A rule adding the packet to a dynamic list.
- Every fifth (and higher) packet from the same source host address within 30 seconds is dropped.
- Every fourth packet from the same source host address within 30 seconds is logged.
- Every fourth packet from the same source host address within 30 seconds is dropped.
- Anything that has not been stopped so far is accepted.

Test your rules, paying attention to the counters with `iptables -vL` before and after connecting via ssh. If you implemented the rules correctly, you should notice that you have to wait a few seconds from the moment you start trying to connect for the fourth time in 30 seconds from the external host.



Visualizing the current list of blocked addresses

The directory `/proc/self/net/xt_recent/` stores each dynamic list you define in a file of the same name as the dynamic list’s identifier. For testing your rules, you can run `cat SSH_COUNTER` to print the contents of the file, e. g., before and after each test to see the changes.

You have now a set of rules that would minimize the impact of an attack on your system but there is still a way for a powerful attacker to carry out a successful attack. Can you speculate how and why?



Milestone

Report your progress to a lab assistant. _____

3.3.9 Building Your Own Firewall

Consider your home network, the services you offer and the service you require from the Internet. Which protocols do you actually use? Which can you block? When you compare these requirements to the rules you just created in this lab, which ones would you keep, which ones not and which additional rules would you add. You may note them down in abbreviated form (which port/protocol, what to do with it, why).

rules to keep:

rules not to keep:

additional rules:



Milestone

Report your progress to a lab assistant. _____

4 nmap: Detecting Server Capabilities

nmap is a free open source utility commonly used for network exploration and security auditing. nmap can determine what hosts are available on the network, what services (application name and version) those hosts are offering, what operating systems (and OS versions) they are running, what type of packet filters/firewalls are in use, and dozens of other characteristics. To view the nmap help, run: `man nmap`



Do not use nmap outside the lab

In this second part of the lab, you will switch from the perspective of protecting a network with a firewall to the other side and explore ways to attack a system using the network mapper nmap. You can safely explore this tool inside this lab environment. Outside the lab, however, remember never to use nmap on a system if you do not have the explicit consent of the system's owner. Even basic scans on one IP address might be considered as attack and have potentially legal implications.

4.1 Enumeration

You will now use nmap to enumerate your outside network from the external host (iptables-□.3). Try to locate the IP address of the unknown server that is connected to the network (see Appendix A). Note that the scan can take up to 5 minutes, so you might want to test your command on a smaller portion of the network before scanning the whole subnet. Also make sure to use options that speed up the scanning, otherwise the scan will take a very long time (up to one hour with some standard settings). Have a look at the `--min-rate` option for that purpose.



Saving time

Since the scan will take some time, pipe the output from nmap to a file and put the process in the background so you can continue your work. Use the `>` operator to pipe the output to a file and the `&` operator to send the process to background:

```
nmap [arguments for nmap] > [output_file] &
```

Which arguments did you use to locate the server?

nmap

What is the server's address, how long time did the scan take and how many addresses did you scan?

server address:

time: number of scanned addresses:

If you get an error message while scanning: Make sure you are root. If you still can't find the server, verify your routing and netmask settings.

4.2 Service Discovery

Now that you have found the server's IP address, try to gather more information, such as the services running on the server. Scan for both TCP and UDP services. Furthermore, try to find out the version number of the discovered services. For the UDP scan you might want to look into the `-F` option.

What arguments did you use for TCP discovery?

nmap

What arguments did you use for UDP discovery?

nmap

UDP discovery is much slower than TCP discovery. Why?

What is the difference between *open*, *closed*, *filtered* and *unfiltered* ports?

open:

closed:

filtered:

unfiltered:

Which services did you find? (If you found about 7 open TCP ports and about 3 open UDP ports you are fine. You might not be able to find version information for all services.)

TCP/UDP	Port	Service	Version

4.3 OS discovery

Finally, try to guess the operation system of the server by using the appropriate nmap operations.

What arguments did you use:

nmap

What operation system did nmap consider to be most likely? (If you do not get any text that makes sense to you, try again using the extra option `--fuzzy`)

What information is used by nmap, to guess the operation system?



Milestone

Report your progress to a lab assistant. _____

5 Cleanup

You do not need to clean up anything, closing the terminal windows is fine. Do NOT shutdown any of the virtual machines (as this can trigger a bug in the hypervisor affecting all other machines).

You are done with the lab, yeah!

6 History

Version	Contribution	Author (Affiliation)	Contact
1.0	First development	Pehr Söderman (ICT/KTH)	pehrs@kth.se
2.0	Adaptation for HT2012	Benjamin Greschbach (CSC/KTH)	bgre@kth.se
3.0	Adaptation for HT2013	Guillermo Rodríguez Cano (CSC/KTH)	gurc@csc.kth.se
3.1	New exercise for HT2013	Guillermo Rodríguez Cano (CSC/KTH)	gurc@csc.kth.se
3.2	Updating for HT2015	Benjamin Greschbach (CSC/KTH)	bgre@csc.kth.se

A Lab Network Map

The network map below shows the connections and information about the addressing of the hosts in the lab network. Remember to replace the placeholders # and □ with your according group values. Furthermore, you have to concatenate the network prefix with the host part to obtain the IP address for a certain interface. The length of the network prefix is given by the network mask³. For example, if your group number # is 42, then `eth0` of the external host should get the address 10.42.0.2 (10.42.0.0/16 being the network prefix and .0.2 the host part), `eth0` of the firewall host should get 192.168.42.1, and so on.

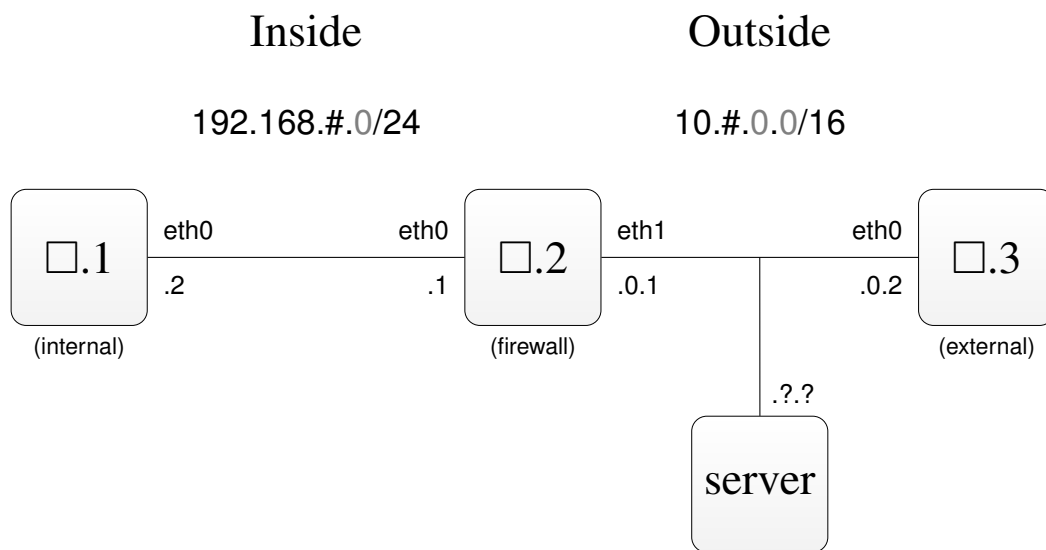


Figure 1: Lab network map.

B FAQ – Common Problems and Fixes

I can't ping between □.1 and □.3 (but between direct neighbors). Check that `route add default gw ADDRESS` was issued **WITHOUT** specifying a netmask after the neighbor's address. You have to take the interface down (`ifconfig eth10 down`) and up again (`ifconfig eth10 up`), after issuing the correct command.

telnet got stuck, how can I exit from it? Issue `Ctrl+C` on the listening netcat `nc` instance. `Ctrl+]`, which is usually used to exit telnet, will in this lab exit the terminal back to the virtual-shell instead.

I can't login with ssh when connecting from one virtual host to another. Make sure to use the username `student` (the only user account on the virtual hosts) when connecting with `ssh`: `ssh student@IPADDRESS`. Then the passphrase `basket field payment same love` will work.

I've set up the correct rules, but ping does not work (no rules match). Note that `ping` uses the ICMP protocol, so it's neither using TCP nor UDP. If your rules match only for TCP and UDP, ping traffic is not affected by them.

³The netmask is written right behind the network address in standard slash-notation: `/24` stands for a netmask of 24 bits (which is equal to 255.255.255.0 in legacy notation). Remember that an IPv4 address has 32 bits, so each number in the dot-notation represents 8 bits.

The virtshell disconnects from an instance when trying the console command on it. Try to write `quit` in the virtshell and reconnect.

The virtshell asks for a password when issuing the boot command. Note that the instance name has to be written case sensitive (capital letter after the dash).

C Flow Diagram of iptables

The following graph shows a simplified and brief scheme of the iptables flow graph. Each chain contains the corresponding tables for the processing of the packets. There are five predefined chains, namely, PREROUTING, FORWARD, INPUT, OUTPUT and POSTROUTING, but not all chains have all the tables (e.g., the INPUT chain only contains the 'mangle' and the 'filter' tables). Note that predefined chains have a default policy (e.g., ACCEPT) which is always applied to the packet traversing the chain when it does not match any of the rules defined for that chain.

Packets start at a given chain, but typically they come either as 'inbound traffic' (that is, from the network card) or from a 'local process', and after traversing the tables of each chain, they will get out of iptables as 'outbound traffic'. Note that depending on the rules the processing of the packets might not be sequential, as the graph shows, and the packets may jump to another chain or even discarded/dropped, but in any case, every packet coming or leaving iptables will traverse one chain at least.

For this lab the only table that you will be using is the one focused on packet filtering, 'filter'.

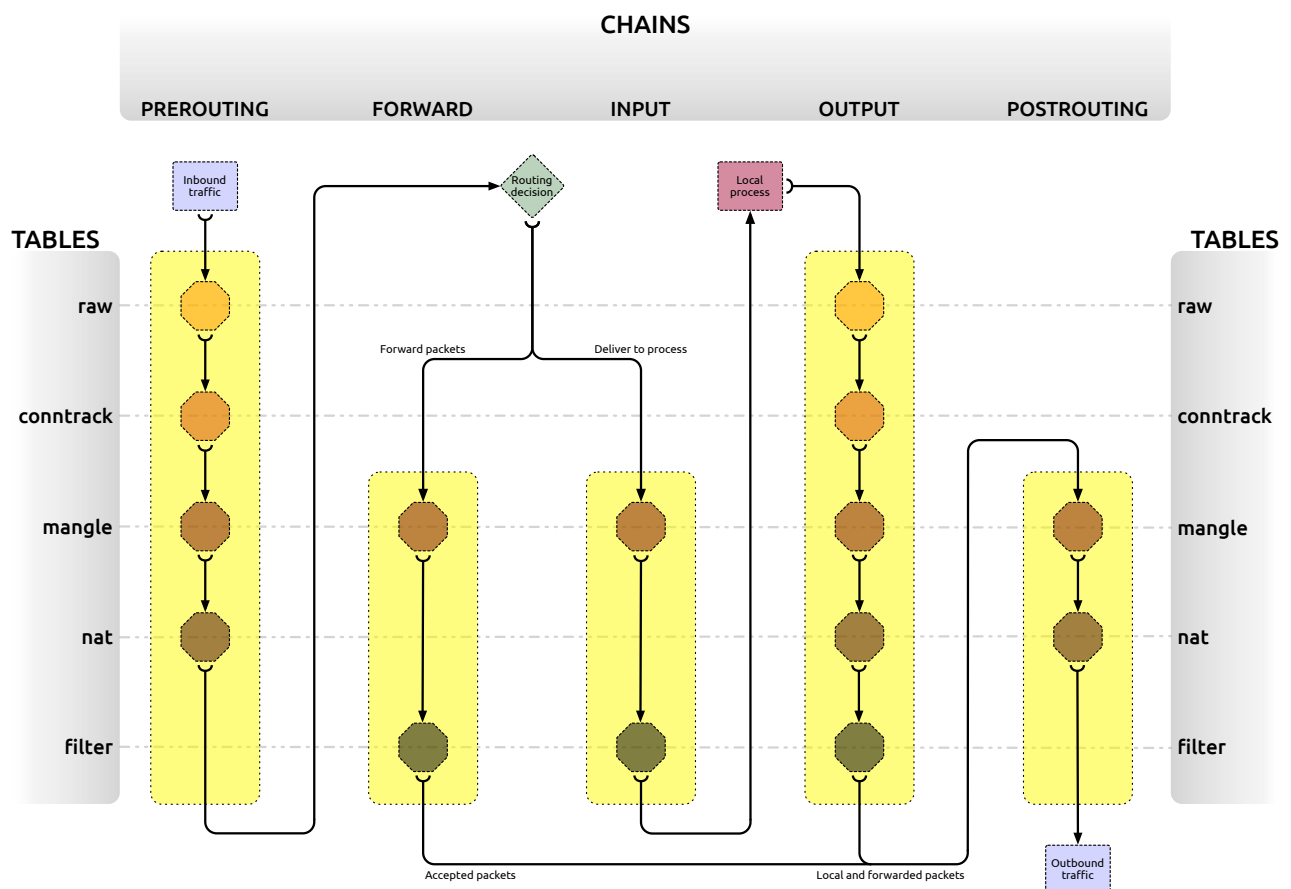


Figure 2: iptables flow graph

D Command Line Basics – in case you are not familiar with the shell

D.1 Manual Pager

You get a reference manual for almost all command line programs by issuing `man PROGRAMNAME`, e. g., `man iptables`. It uses the same shortcuts as the file viewer `less`, which provides the following actions: Pressing `Q` will exit the viewer (quit). To scroll up or down half a page press `Ctrl+U` or `Ctrl+D` (or `Space/Shift+Space` for whole page scrolling). `/keyword` starts a search for "keyword" (only downwards from current position). Pressing `N` or `Shift+N` will get you to the next or previous occurrence of the keyword. `G` or `Shift+G` will goto the beginning or end of the file.

D.2 Viewing a File

Use `ls` to list the files in the current directory (`cd ..` and `cd DIRECTORY` to change the directory). For viewing a file there are several possibilities, `less FILENAME` being one of the more convenient ones (see Section D.1 above for navigation shortcuts). If you want to read a file continuously (because you expect data to be written to it by another program while reading it) you can use `tailf FILENAME`, which you have to exit by pressing `Ctrl+C`.

D.3 Editing a File

There are several file editors you can use on the command line, but if you are not familiar with `vim` or `emacs`, the best choice is probably to use `nano`. You open a file for editing by calling `nano FILENAME`. Then you can edit the file and move around with arrow keys. All available commands are displayed at the bottom, where `^` denotes the `Ctrl` key, so you can save the file ("WriteOut") with `Ctrl+O` and exit `nano` with `Ctrl+X`.

D.4 Background Jobs

You can use the `&` character after any command that may take more time, to start it running in background (so that you can continue to use the terminal while the command is executed).

For example `nmap -sP 10.0.0.0/16 > scanresults.txt &`

Alternatively, you can press `Ctrl+Z` (suspend), which stops the current job and sends it to the background. To see all jobs were sent to the background, run `jobs` which also prints out a number for each job. `bg JOBNUMBER` or `fg JOBNUMBER` can be used to either continue a job in the background or foreground. Pressing `Ctrl+C` while a job is running in foreground will kill (exit) it.

D.5 Clipboard – Copy+Paste

Note that `Ctrl+C` on the command line is used to kill the current program. Most terminals support `Ctrl+Shift+C` (after selecting lines with the mouse) and `Ctrl+Shift+V` as shortcuts for copy and paste to and from the system clipboard. Using the middle mouse button to paste selected text is also supported frequently.